

Step 1: The Ideation Hopper (Problem Identification)

- **The UI:** A simple, repeatable card layout where they can click "Add Problem" up to 5 times.
- **The Inputs:**
 - **Problem Title:** (e.g., "Manual Loan Application Review")
 - **The Business Challenge:** What is broken right now? (e.g., "Takes 3-5 days, 35% drop-off")
 - **End Users:** Who suffers from this?
 - **Current Non-AI Solution:** How is it solved today?

Step 2: The "AI vs. Automation" Sieve

- **The UI:** The 5 problem cards from Step 1 are carried over. Beside each, they have a checklist of the core base questions you teach.
- **The Inputs:**
 - *Are the rules 100% definable? (Yes/No)*
 - *Does this require pattern recognition, prediction, or natural language generation? (Yes/No)*
 - *Could a traditional script achieve this with lower risk? (Yes/No)*
- **The Action:** Based on their answers, the form forces them to "Discard" the purely automation tasks and "Select" their top 2-3 AI-suited problems to move forward.

Step 3: The Feasibility Matrix (Deep Dive)

- **The UI:** Side-by-side comparison columns for their surviving 2-3 problems. This forces them to confront the reality of building the system.
- **The Inputs:**
 - **Data Readiness:** What is the volume? Are there labels? If the real dataset is too sparse, imbalanced, or blocked by privacy laws, is synthetic data generation a viable path?
 - **Technical Readiness:** What is the compute budget? What modeling libraries (e.g., PyTorch, Scikit-learn) or architectures are required?
 - **Regulatory & Compliance:** Are there specific constraints? (e.g., EU AI Act, GDPR, or specific domain rules like FINMA for finance or HL7 FHIR for healthcare).

Step 4: The Final Specification Generator

- **The UI:** They select the *single most feasible problem* from Step 3. The form now expands into a detailed page that maps exactly to the headers of your final briefing document. It auto-fills what they've already typed and asks them to flesh out the rest.

- **The Inputs (Mapping to the Document):**
 - **Executive Summary & Objectives:** Auto-populated from Step 1, but they are asked to refine it and add strict *Success Metrics* (e.g., "Reduce decision time to under 30 mins").
 - **Data Assets:** Pulled from Step 3, asking for specific formats (e.g., "1.2M records in CSV").
 - **Regulatory & Compliance Requirements:** Expanding on Step 3, asking them to explicitly state the *Model Development & Testing* constraints and the *Deployment & Operations* safeguards (like human-in-the-loop requirements).
 - **Stakeholders:** A dynamic table where they add rows for each stakeholder (e.g., Chief Risk Officer), their Primary Concerns (e.g., "Model accuracy"), and Success Metrics (e.g., "Default rates").
 - **Solution Components:** A checklist or text area to define the actual pipeline steps (e.g., "Text Processing Pipeline," "Bias Detection System").

Once they hit "Submit," the web app stitches these fields together into that clean, 2-page PDF/Markdown format.